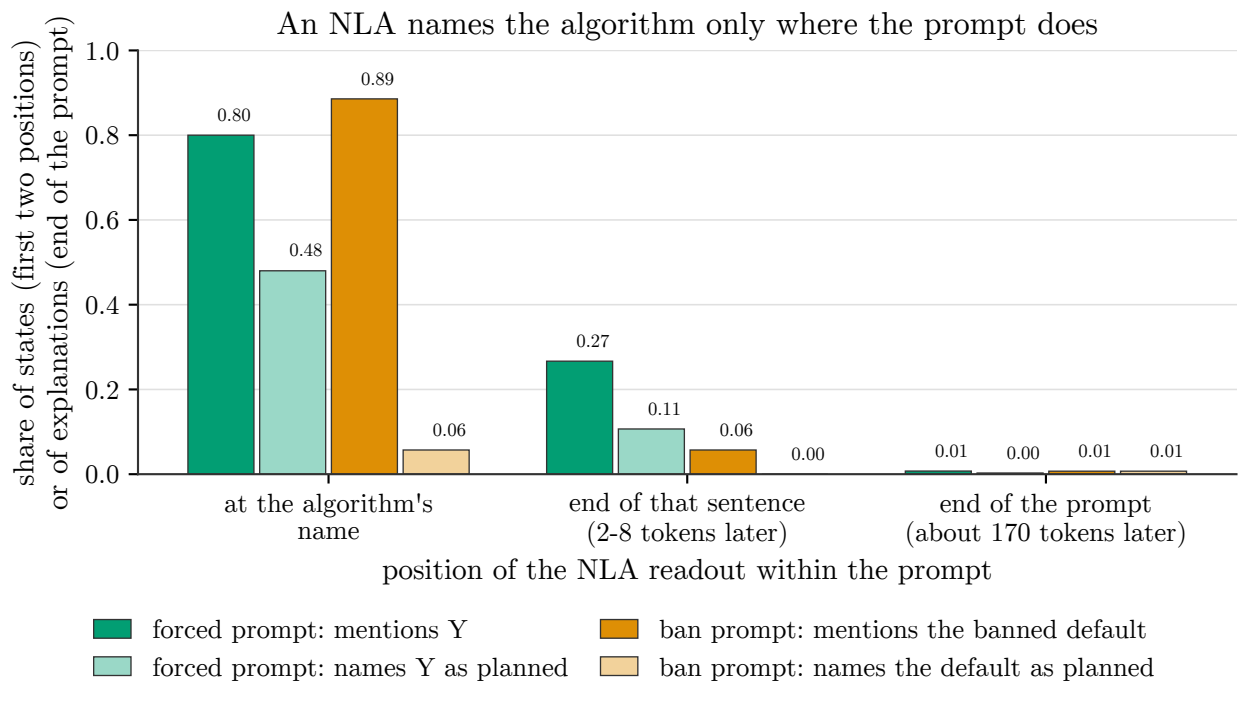
An NLA names the algorithm only where the prompt does
share of states (first two positions)
or of explanations (end of the prompt)
1.0
0.8
0.6
0.4
0.2
0.0
0.80
0.48
0.89
0.06
0.27
0.11
0.06
0.00
0.01
0.00
0.01
0.01
at the algorithm's name
end of that sentence (2-8 tokens later)
end of the prompt (about 170 tokens later)
position of the NLA readout within the prompt
forced prompt: mentions Y
forced prompt: names Y as planned
ban prompt: mentions the banned default
ban prompt: names the default as planned